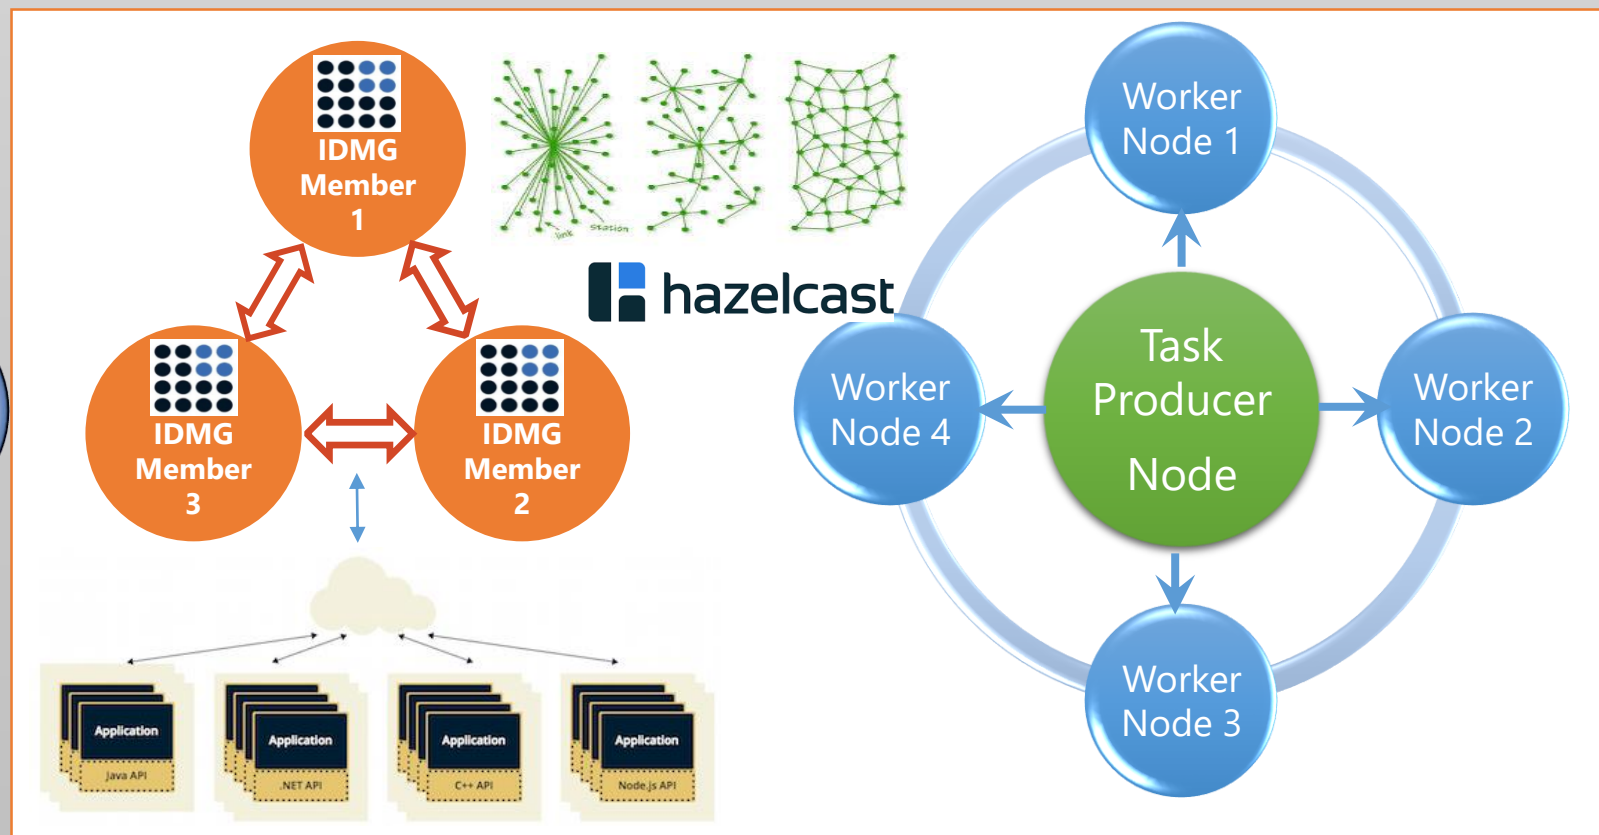
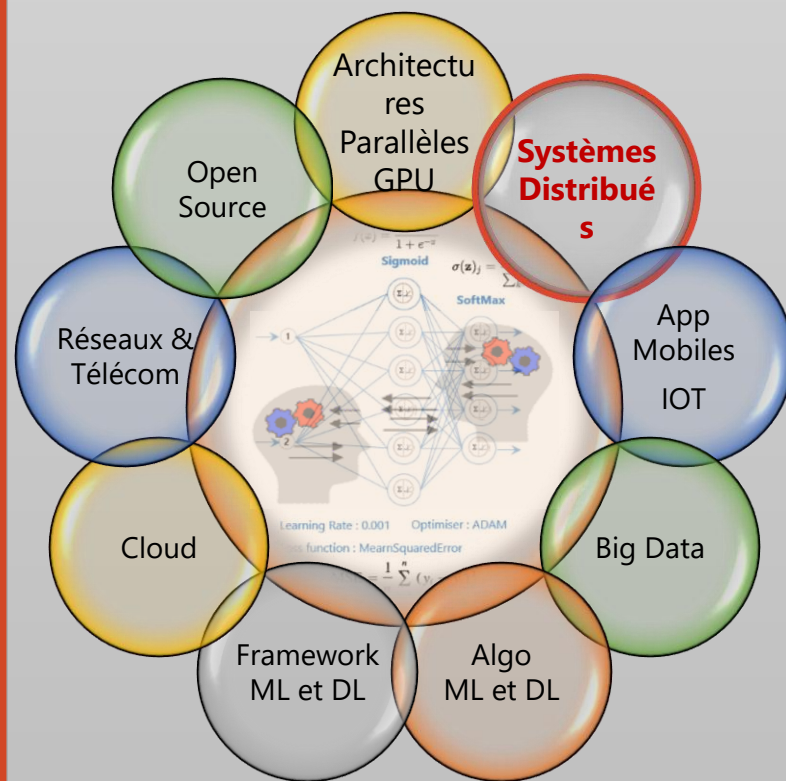


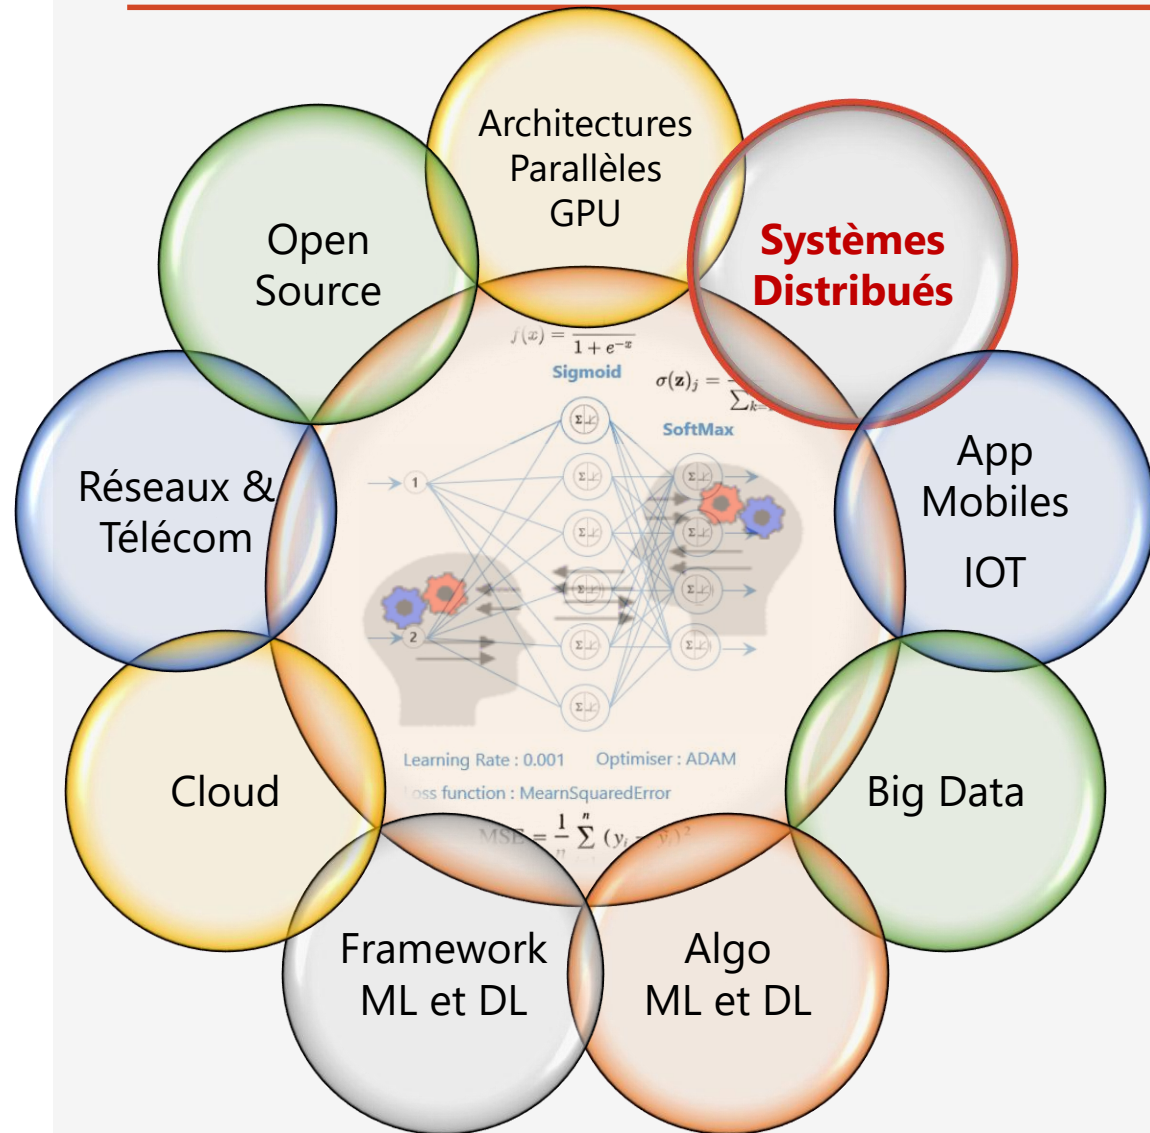
# Distributed Computing and Caching with Hazelcast EcoSystem

IDMG : In Distributed Memory Grid

By : Mohamed YOUSSEFI  
Lab. SSDIA, ENSET Mohammeda,  
Hassan II University of Casablanca  
Morocco  
[med@youssefi.net](mailto:med@youssefi.net)



# Systemes Distribués



## ➤ Systèmes Distribués

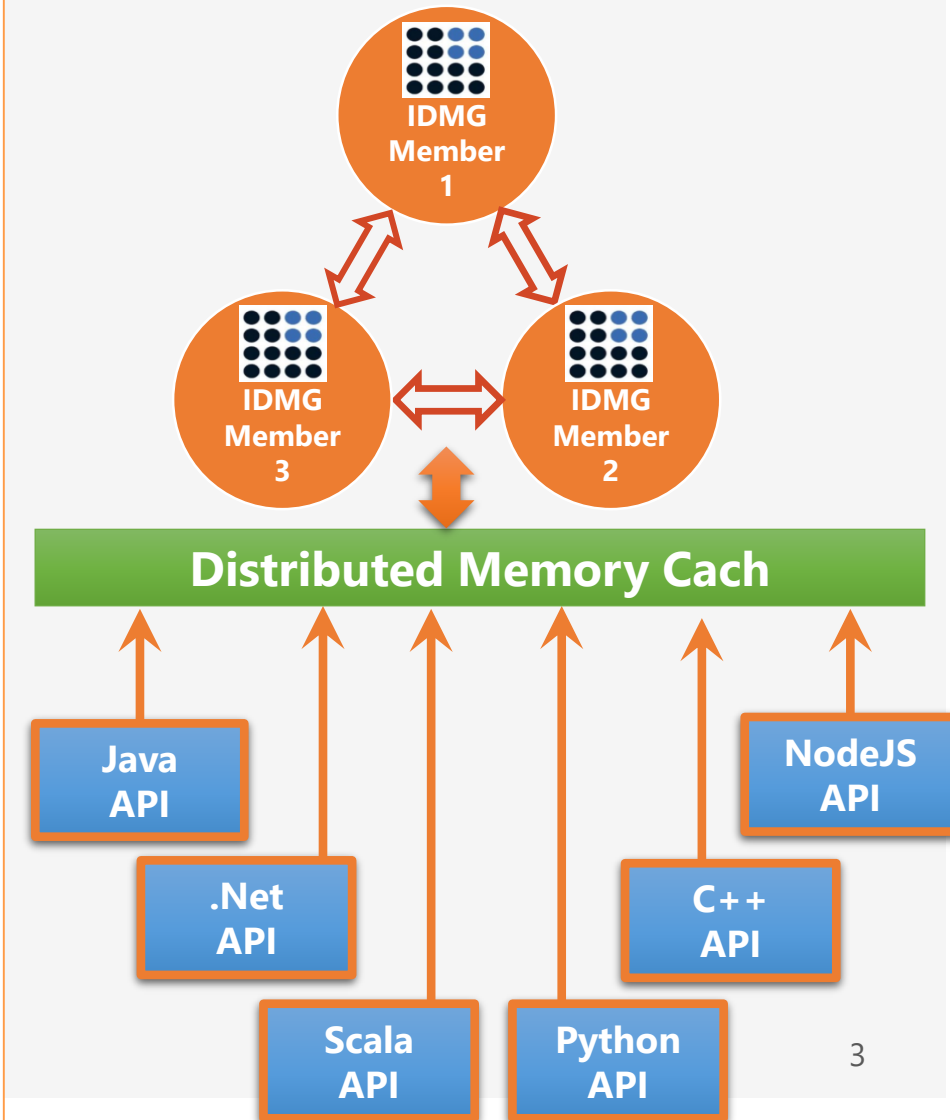
- Middlewares RMI, CORBA, JMS, MPI
- Grilles de Clacul
- Protocoles de messageries asynchrones:
- **AMQP, MQTT, STOMP**
- Brokers : **RabbitMQ, ActiveMQ**
- Caches mémoires Distribués : **Hazelcast**
- Middlewares **SMA** : **JADE**

## Systèmes Distribués Middlewares



# C'est quoi Hazelcast ?

- Hazelcast IMDG est un middleware Open source en Java, qui permet de créer un cache mémoire distribué.
- Dans une grille Hazelcast, les données sont réparties uniformément entre les nœuds d'un groupe d'ordinateurs, ce qui permet
  - Un stockage distribué Scalable (Distributed Memory Cache)
  - Un traitement distribué Scalable (Distributed Computing).
  - Réplication des données sur plusieurs nœud pour tolérance aux pannes
- Ces techniques réduisent la charge de requête sur les bases de données et améliorent les performances des systèmes distribués.



# Utilisations de Hazelcast

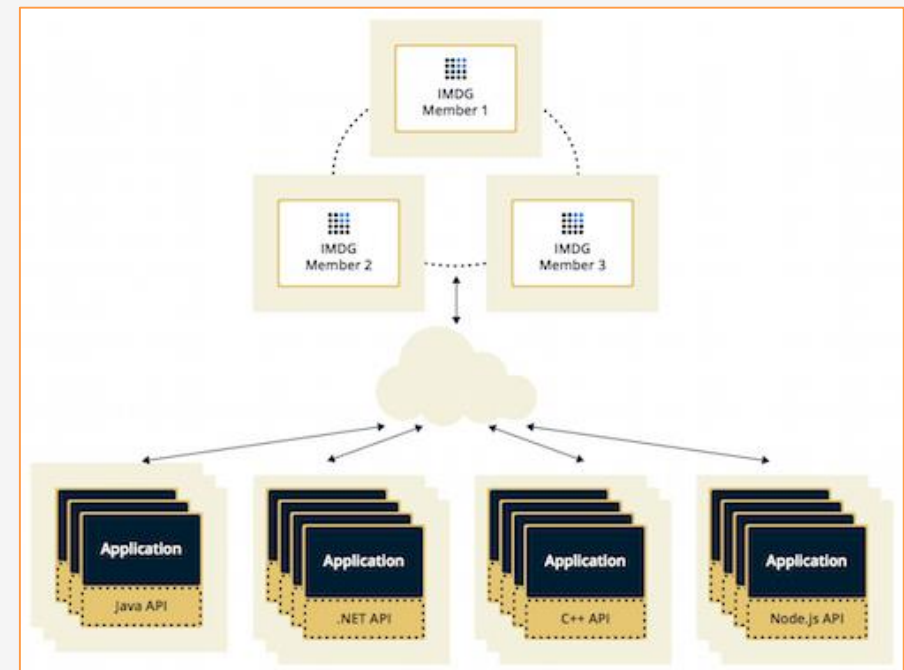
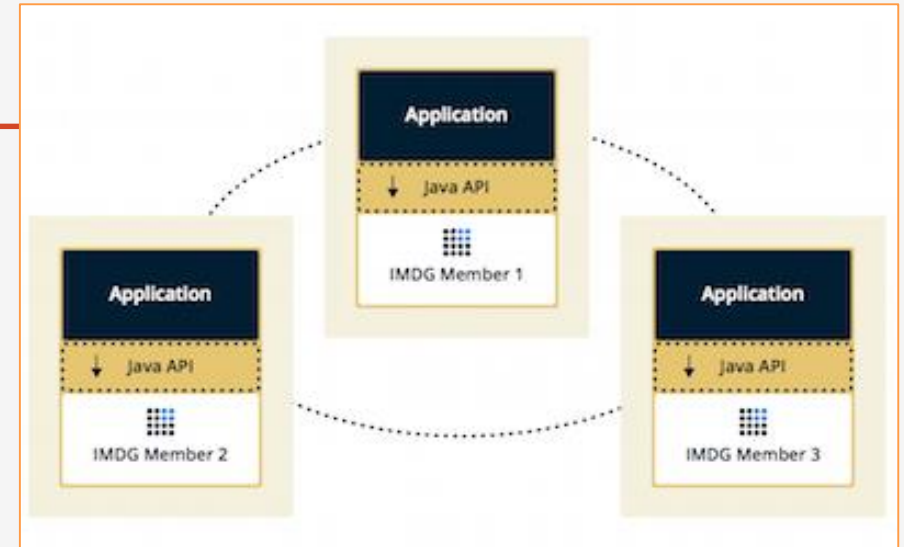
---

Hazelcast est une solution quand vous avez besoin de:

- Applications analytiques nécessitant un traitement de données volumineuses en partitionnant les données (Big Data)
- Pour conserver les données fréquemment consultées dans la grille mémoires
- Un cache mémoire distribué hautement disponible pour les applications
- Un magasin de données principal pour les applications avec des exigences de performances, d'évolutivité et de latence maximales
- In Memory NoSQL data base de type Clé-Valeurs
- Solution de messagerie (publier / souscrire) à très rapide avec un scalabilité entre les applications
- Solution faire distribuer les traitements (Distributed Computing)
- Une alternative aux autres solutions comme Coherence and Terracotta.

# C'est quoi Hazelcast ?

- Hazelcast est implémenté en Java et possède des clients pour Java, C / C ++, .NET, REST, Python, Go et Node.js.
- Contrairement à beaucoup de solutions NoSQL, Hazelcast est peer-to-peer. Il n'y a pas de maître et d'esclave;
- il n'y a pas point de défaillance unique (SPOF).
- Système d'équilibrage de charges : Tous les membres stockent des quantités égales de données et font des quantités égales de En traitement.
- On peut intégrer Hazelcast dans une application existante
- ou l'utiliser pour faire de votre application est un client pour un cluster Hazelcast.





# Caractéristiques de Hazelcast


---

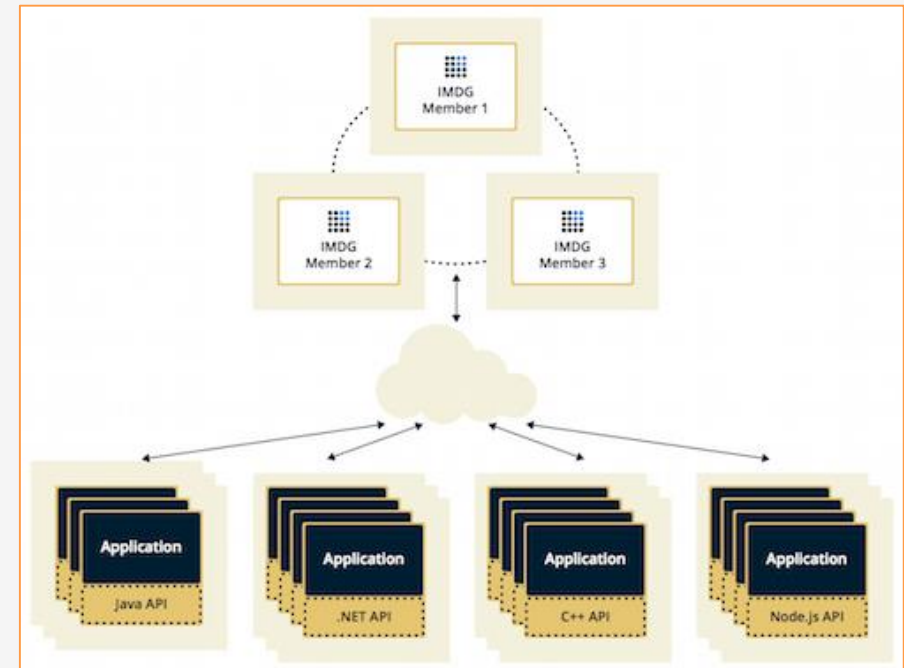

- Hazelcast est une solution Opensource
- Pour utiliser Hazelcast vous avez besoin uniquement d'un seul fichier JAR.
- Hazelcast fournit une implementation distribuée des structures de données classiques acomme
  - Map, Queue, MultiMap, Topic, Lock Executor.
- On peut ajouter de Nouvelles implementation de structures de données distribuées en utilisant le Service Programming Interface (SPI)
- L'architecture de Hazecast est comlètement distribuée :
  - Il n'y a pas de noeud Master centralisé
  - Pas de point d'échec (SPOF)
- Tous les noeuds sont configurés pour être identiques
- Quand les capacités mémoire et de calcul nécessitent de croitre, Il suffit de démarer de nouveaux membres faisant partie du Cluster (Scalability)
- Les données sont résilientes à l'échec d'un membre , vu que des backups des données sont duppliquées dans plusieurs noeuds du cluster

# Intégration de Hazel Cast

```
<dependencies>
  <dependency>
    <groupId>com.hazelcast</groupId>
    <artifactId>hazelcast</artifactId>
    <version>3.12</version>
  </dependency>
  <dependency>
    <groupId>com.hazelcast</groupId>
    <artifactId>hazelcast-client</artifactId>
    <version>3.12</version>
  </dependency>
</dependencies>
```

 hazelcast-3.12.jar

 hazelcast-client-3.12.jar



# Première Application Hazelcast

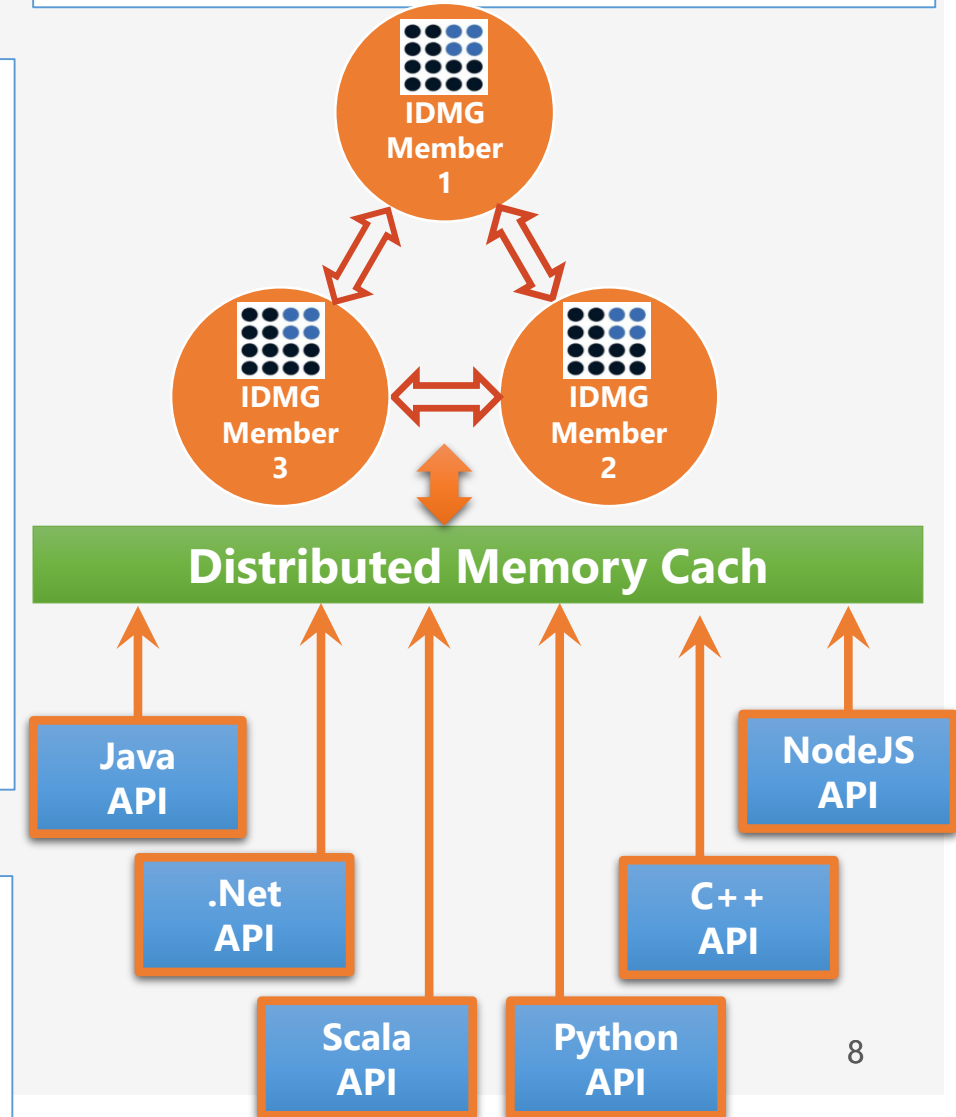
## Démarrage d'un nœud du cluster hazelcast

```
Config cfg = new Config();
HazelcastInstance instance = Hazelcast.newHazelcastInstance(cfg);
Map<Integer, String> mapCustomers = instance.getMap("customers");
mapCustomers.put(1, "Joe");
mapCustomers.put(2, "Ali");
mapCustomers.put(3, "Avi");
System.out.println("Customer with key 1: " + mapCustomers.get(1));
System.out.println("Map Size:" + mapCustomers.size());
Queue<String> queueCustomers = instance.getQueue("customers");
queueCustomers.offer("Tom");
queueCustomers.offer("Mary");
queueCustomers.offer("Jane");
System.out.println("First customer: " + queueCustomers.poll());
System.out.println("Second customer: " + queueCustomers.peek());
System.out.println("Queue size: " + queueCustomers.size());
```

## Démarrage d'un client du cluster hazelcast

```
ClientConfig clientConfig = new ClientConfig();
HazelcastClient client = HazelcastClient.newHazelcastClient(
    clientConfig);
IMap map = client.getMap( "customers" );
System.out.println( "Map Size:" + map.size() );
```

- En exécutant cette application en plusieurs instances, vous créez une grille mémoire distribuée
- Les données des deux collections Map et Queue sont réparties sur plusieurs nœuds.
- Les deux collections distribuées sont partagées à toutes les instances.



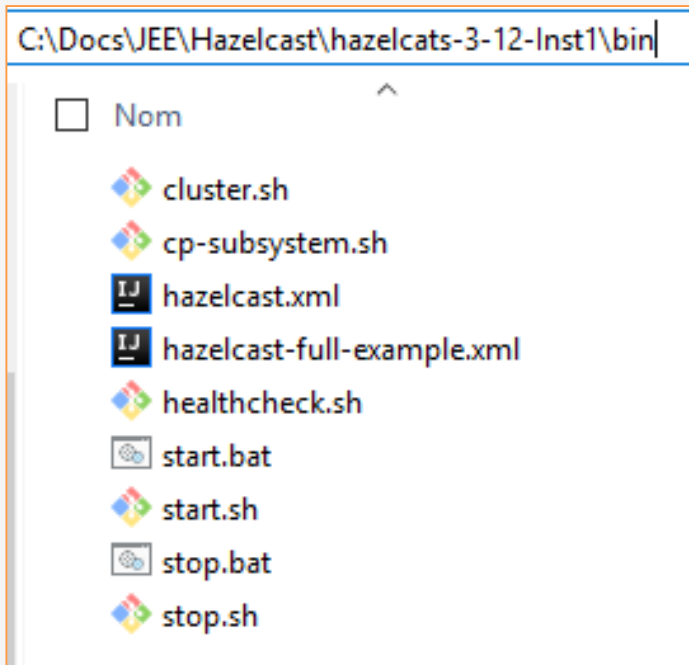


# Démarrage des nœuds Hazelcast avec les scripts fournis

- Télécharger la distribution de Hazelcast DMG:
  - <https://hazelcast.org/download/>

## Démarrage d'un nœud Hazelcast

```
Invite de commandes
C:\Docs\JEE\Hazelcast\hazelcats-3-12-Inst1\bin>start.bat
#####
# RUN_JAVA=C:\Program Files\Java\jdk1.8.0_151\bin\java
# JAVA_OPTS=
# starting now...."
#####
C:\Docs\JEE\Hazelcast\hazelcats-3-12-Inst1\bin>
```



```
Members {size:1, ver:1} [
  Member [192.168.1.49]:5701 - e19e9197-53a4-4930-84c4-d4bf6a213473 this
]
```

## Démarrage d'un autre nœud Hazelcast

```
Invite de commandes
C:\Docs\JEE\Hazelcast\hazelcats-3-12-Inst2\bin>start.bat
#####
# RUN_JAVA=C:\Program Files\Java\jdk1.8.0_151\bin\java
# JAVA_OPTS=
# starting now...."
#####
C:\Docs\JEE\Hazelcast\hazelcats-3-12-Inst2\bin>
```

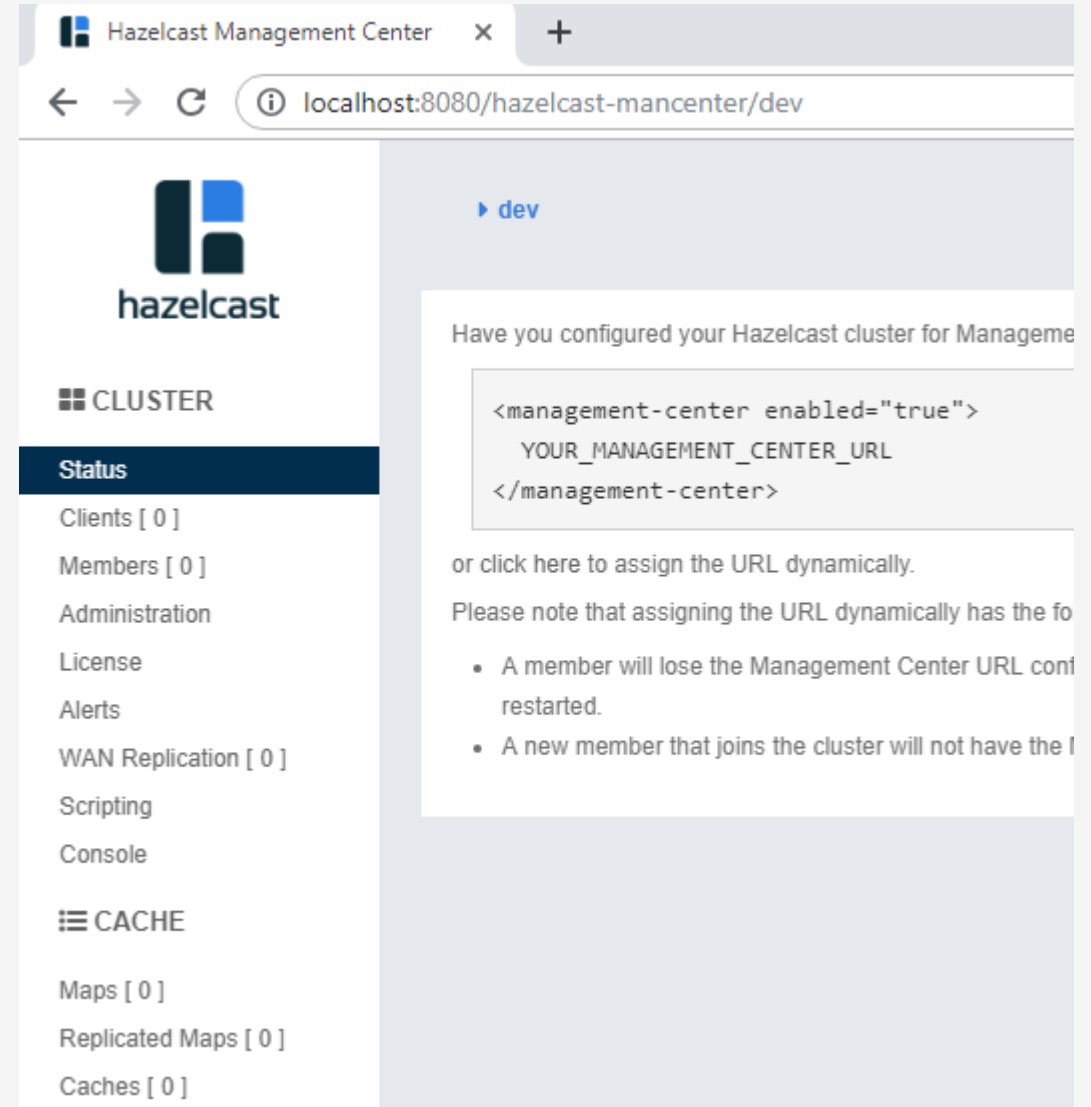
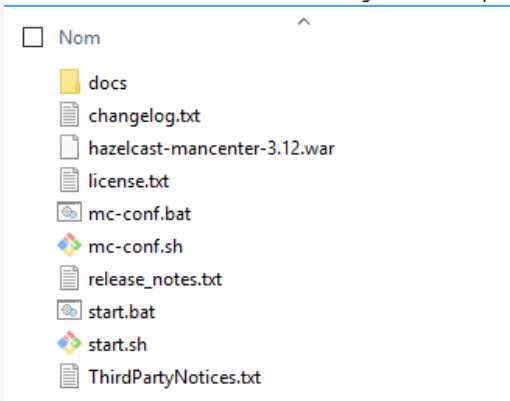
```
Members {size:2, ver:2} [
  Member [192.168.1.49]:5701 - e19e9197-53a4-4930-84c4-d4bf6a213473
  Member [192.168.1.49]:5702 - 816f2c72-fa59-4119-8839-1dcc03f710dc this
]
```

# Démarrage de la console de monitoring et de management Hazelcast


## Démarrage de la console de management

```
Invite de commandes - start.bat
C:\Docs\JEE\Hazelcast\hazelcast-3.12\management-center>start.bat
2019-04-09 11:25:04 [main] INFO c.h.webmonitor.config.BuildInfo - Management Center 3.12
2019-04-09 11:25:08 [main] INFO c.h.w.storage.DiskUsageMonitor - Monitoring C:\Users\med\hazelcast-mc [mode=purge, interval=1000ms, limit=512 MB]
2019-04-09 11:25:17 [main] INFO c.h.w.s.s.impl.DisableLoginStrategy - Login will be disabled for 5 seconds after 3 failed login attempts. For every 3 consecutive failed login attempts, disable period will be multiplied by 10.
2019-04-09 11:25:17 [main] INFO c.h.webmonitor.config.SqlDbConfig - Checking DB for required migrations.
2019-04-09 11:25:18 [main] INFO c.h.webmonitor.config.SqlDbConfig - Number of applied DB migrations: 0.
Hazelcast Management Center successfully started at http://localhost:8080/hazelcast-mancenter
```

C:\Docs\JEE\Hazelcast\hazelcast-3.12\management-center\



# Démarrage de la console de monitoring et de management Hazelcast

  
hazelcast

CLUSTER

Status

- Clients [ 0 ]
- Members [ 2 ]
- Administration
- License
- Alerts
- WAN Replication [ 0 ]
- Scripting
- Console

CACHE





- Maps [ 0 ]
- Replicated Maps [ 0 ]
- Caches [ 0 ]
- MultiMaps [ 0 ]
- PN Counters [ 0 ]
- ID Generators [ 0 ]

MESSAGING


- Queues [ 0 ]
- Topics [ 0 ]
- Reliable Topics [ 0 ]

devDocumentationTime Traveladmin Last Login: 09/04/2019 à 09:09:24devLog Out

Memory Utilization

Member	Used Heap ↕	Total Heap ↕	Max. Heap ↕	Heap Usage Percent...	Used Heap	Native Memory Max ↕	Native Memory Used ↕	Native Memory Free ↕	Native Memory Used	GC Major Count ↕	GC Major Time ↕	GC Minor Count ↕	GC Minor Time ↕
192.168.1.49:57000	40.05 MB	104.50 MB	1.75 GB	2,23 %		0	0	0		1	28.00ms	4	20.00ms
192.168.1.49:57001	38.51 MB	104.50 MB	1.75 GB	2,15 %		0	0	0		1	28.00ms	4	19.00ms

Heap Memory Distribution



mapotherfree

Partition Distribution

No data

Cluster State

CPU Utilization

# Connecter un nœud hazelcast à hazelcast management center

```
ManagementCenterConfig mcfg =new
ManagementCenterConfig();

mcfg.setEnabled(true);

mcfg.setUrl("http://localhost:8080/hazelc
ast-mancenter");

Config cfg = new Config();

cfg.setManagementCenterConfig(mcfg);

HazelcastInstance instance =
Hazelcast.newHazelcastInstance(cfg);
```

The screenshot shows the Hazelcast Management Center interface. The left sidebar contains a navigation menu with sections: CLUSTER, CACHE, and MESSAGING. The 'CLUSTER' section is expanded, showing sub-items: Status, Clients [0], Members [3], Administration, License, Alerts, WAN Replication [0], and Scripting. The 'Console' item is selected. The main area displays a terminal window with the following content:

```
localhost:8080/hazelcast-mancenter/dev/console

192.168.1.49:5703 [default]$ customers.values
type 'help' for help
192.168.1.49:5703 [default]$ customets.keys
type 'help' for help
192.168.1.49:5703 [default]$ ns cutomers
namespace: customers
192.168.1.49:5703 [cutomers]$ m.values
Total 0
192.168.1.49:5703 [cutomers]$ m.size
Size: 0
192.168.1.49:5703 [cutomers]$ ns customers
namespace: customers
192.168.1.49:5703 [customers]$ m.size
Size: 3
192.168.1.49:5703 [customers]$ m.values
Avi
Ali
Joe
Total 3
192.168.1.49:5703 [customers]$
```

# Connecter un nœud hazelcast à hazelcast management center

```
ManagementCenterConfig mcfg =new ManagementCenterConfig();
mcfg.setEnabled(true);
mcfg.setUrl("http://localhost:8080/hazelcast-mancenter");
Config cfg = new Config();
cfg.setManagementCenterConfig(mcfg);
HazelcastInstance instance = Hazelcast.newHazelcastInstance(cfg);
Map<Integer, String> mapCustomers = instance.getMap("customers");
for (int i = 0; i <10000 ; i++) {
    mapCustomers.put(i, "Data "+i);
}
```

Après arrêt d'une instance

Members	Entries
192.168.1.49:...	4984
192.168.1.49:...	5016
TOTAL	10000

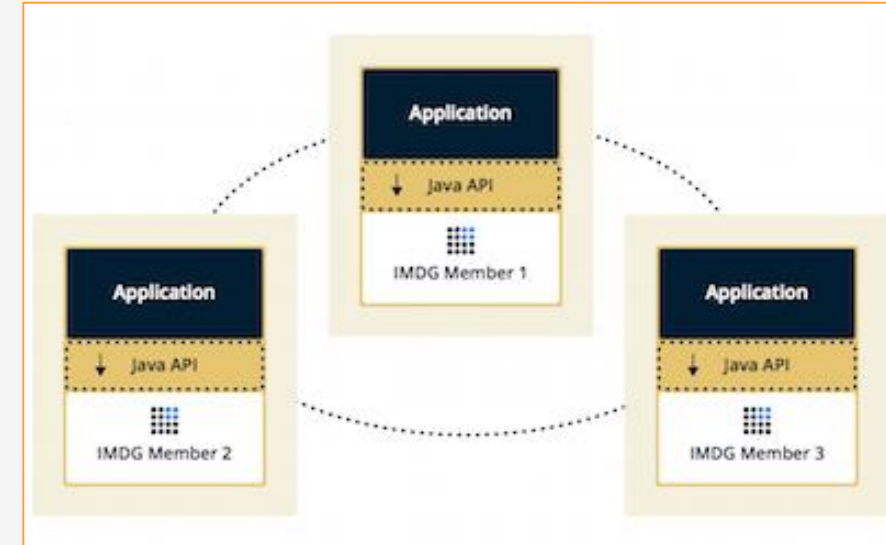
3 instances

Members	Entries	Gets	Puts	Removals	Entry Memor...	Backups	Backup Mem...	Events	Hits	
192.168.1.49:...	3354	0	10000	0	474.57 kB	3354	474.57 kB	0	3354	0
192.168.1.49:...	3317	0	10000	0	469.33 kB	3347	473.58 kB	0	3317	0
192.168.1.49:...	3329	0	0	0	471.03 kB	3299	466.79 kB	0	3329	0
TOTAL	10000	0	20000	0	1.38 MB	10000	1.38 MB	0	10000	0



# Distributed Computing with Hazelcast : Callable Task


- Hazelcast offers IExecutorService for you to use in distributed environments.
- It implements `java.util.concurrent.ExecutorService` to serve the applications requiring computational and data processing power.
- With IExecutorService, you can execute tasks asynchronously and perform other useful tasks.
- If your task execution takes longer than expected, you can cancel the task execution.
- Tasks should be Serializable since they are distributed.
- In the Java Executor framework, you implement tasks two ways: Callable or Runnable.
  - Callable: If you need to return a value and submit it to Executor, implement the task as **`java.util.concurrent.Callable`**.
  - Runnable: If you do not need to return a value, implement the task as **`java.util.concurrent.Runnable`**.



# Data Partition

- Les fragments Hazelcast sont appelés des **partitions**.
- Les partitions sont des segments de mémoire pouvant contenir des centaines, voire des milliers, d'entrées de données, en fonction de la capacité de mémoire de votre système.
- Chaque partition Hazelcast peut avoir **plusieurs répliques (Copies)**, qui sont répartis entre les membres du cluster.
- Une des répliques devient la **réplique principale** et d'autres sont appelés des **sauvegardes**.
- Un membre du cluster qui possède le réplica principal d'une partition est appelé **partition propriétaire**.
- Lorsque vous lisez ou écrivez une entrée de données particulière, vous **parlez de manière transparente au propriétaire de partition** contenant les données.
- Par défaut, Hazelcast propose **271 partitions**.
- Lorsque vous démarrez un cluster avec un seul membre, il possède toutes les 271 partitions.
- En démarrant d'autres membres **les partitions sont distribuées équitablement aux membres du cluster**


## Avec un membre



Node

## Avec deux membre

- En noir, les partions primaires
- En bleu, les partitions backup



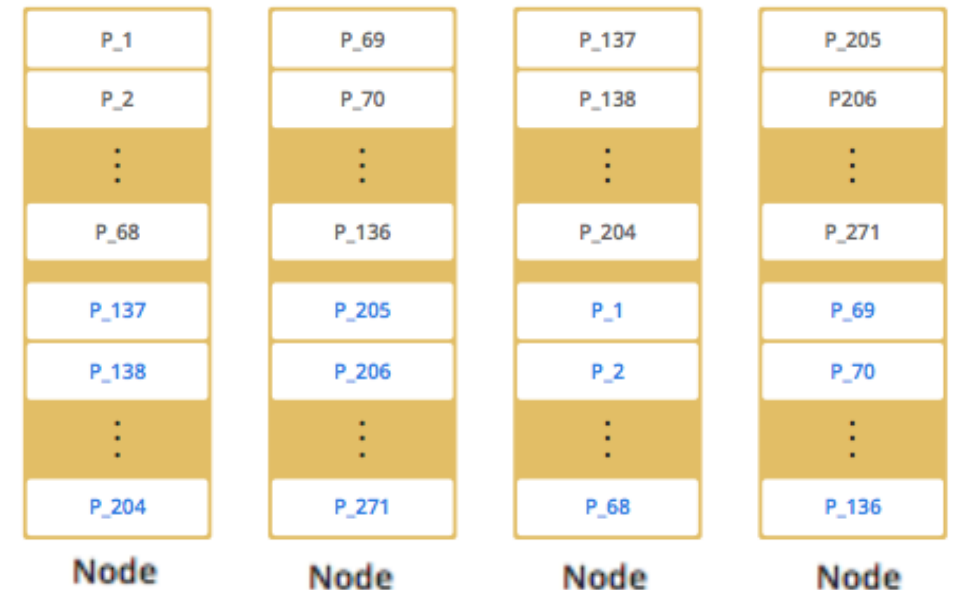
Node

Node

# Data Partition

- Les fragments Hazelcast sont appelés des partitions.
- Les partitions sont des segments de mémoire pouvant contenir des centaines, voire des milliers, d'entrées de données, en fonction de la capacité de mémoire de votre système.
- Chaque partition Hazelcast peut avoir plusieurs répliques (Copies), qui sont répartis entre les membres du cluster.
- Une des répliques devient la réplique principale et d'autres sont appelés des sauvegardes.
- Un membre du cluster qui possède le réplica principal d'une partition est appelé partition propriétaire.
- Lorsque vous lisez ou écrivez une entrée de données particulière, vous parlez de manière transparente au propriétaire de partition contenant les données.
- Par défaut, Hazelcast propose 271 partitions.
- Lorsque vous démarrez un cluster avec un seul membre, il possède toutes les 271 partitions.
- En démarrant d'autres membres les partitions sont distribuées équitablement aux membres du cluster

## Pour un cluster de 4 membre



- Hazelcast distribue les partitions principale et secondaire (backup) de manière égale entre les membres du grappe. Les répliques de sauvegarde des partitions sont conservées pour la redondance.

# Hazelcast Configuration

- Vous pouvez configurer Hazelcast en utilisant un ou plusieurs combinaisons des options suivantes:
  - manière déclarative (XML or a YAML File)
  - manière programmatique (Code Java ou autre)
  - Utilisation des propriétés du système Hazelcast
  - Dans le contexte du Framework Spring
  - Ajout dynamique de la configuration sur un cluster en cours d'exécution

## Configuration programmatique

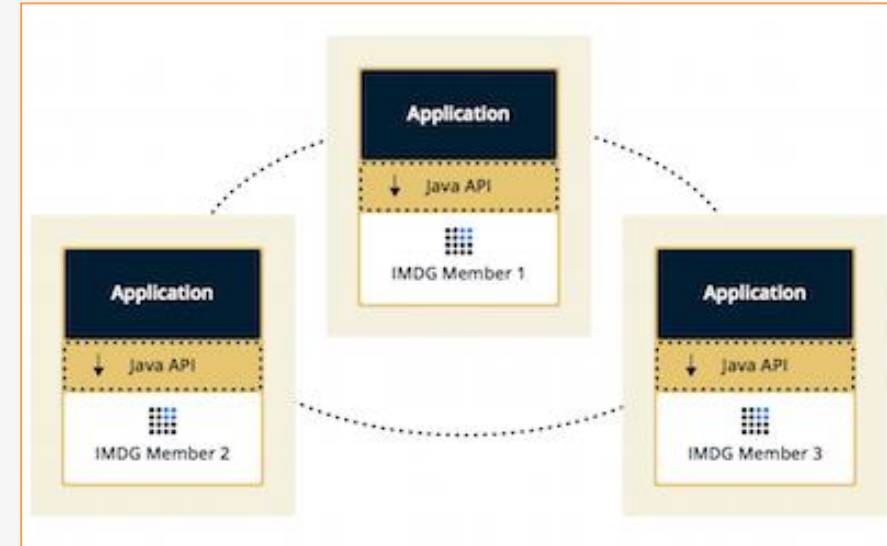
```
Config config = new Config();
config.getNetworkConfig().setPort( 5900 )
    .setPortAutoIncrement( false );
MapConfig mapConfig = new MapConfig();
mapConfig.setName( "testMap" )
    .setBackupCount( 2 )
    .setTimeToLiveSeconds( 300 );
HazelcastInstance member = Hazelcast.newHazelcastInstance(config );
```

## Configuration déclarative : hazelcast.xml

```
<hazelcast>
  <group>
    <name>dev</name>
  </group>
  <management-center enabled="false">http://localhost:8080/mancenter</management-center>
  <network>
    <port auto-increment="true" port-count="100">5701</port>
    <join>
      <multicast enabled="true">
        <multicast-group>224.2.2.3</multicast-group>
        <multicast-port>54327</multicast-port>
      </multicast>
      <tcp-ip enabled="false">
        <interface>127.0.0.1</interface>
        <member-list>
          <member>127.0.0.1</member>
        </member-list>
      </tcp-ip>
    </join>
  </network>
  <map name="default">
    <time-to-live-seconds>0</time-to-live-seconds>
  </map>
</hazelcast>
```

# Discovery Mechanisms

- Un cluster Hazelcast est un réseau de membres qui exécutent Hazelcast.
- Un membre peut se joindre automatiquement au cluster pour former un cluster.
- Cette jonction automatique a lieu avec divers mécanismes de découverte que les membres du cluster utilisent pour se retrouver.
  - TCP
  - Multicast (UDP)
  - Cloud Discovery : (AWS, GCP, JClouds, Azure, Zookeeper, PCF, OpenShift, Eureka, Kubernetes, etc. )
- Notez que:
  - Par défaut Hazelcast utilise le mode Multicast (Ce qui est pratique mais déconseillé en production)
  - Dans un cluster, les communications entre ses membres se font toujours via TCP / IP, quel que soit le mécanisme de découverte utilisé.





# Data Structures

---

- Hazelcast has two types of distributed objects in terms of their partitioning strategies:
  1. Data structures where each partition stores a part of the instance, namely partitioned data structures. :
    - **Map**
    - **MultiMap**
    - Cache (Hazelcast JCache implementation)
    - Event Journal
  2. Data structures where a single partition stores the whole instance, namely non-partitioned data structures :
    - **Queue, Set, List, Ringbuffer**
    - Lock, Isemaphore, IAtomicLong, IAtomicReference
    - FlakeldGenerator, ICountdownLatch, Cardinality Estimator, PN Counter

# Data Structures : Exemple d'utilisation de Map et Queue

---

```
HazelcastInstance hzInstance = Hazelcast.newHazelcastInstance();
```

## Distributed Map

```
Map<String, String> capitalcities = hzInstance.getMap( "capitals" );  
capitalcities.put( "1", "Rabat" );  
capitalcities.put( "2", "Paris" );
```

## Distributed Queue

```
BlockingQueue<MyTask> queue = hzInstance.getQueue( "tasks" );  
queue.put( new MyTask() );  
MyTask task = queue.take();  
boolean offered = queue.offer( new MyTask(), 10, TimeUnit.SECONDS );  
task = queue.poll( 5, TimeUnit.SECONDS );  
if ( task != null ) {  
    //process task  
}
```

# Data Structures : Queue (TaskProducer=>TaskConsumer)

```
public class ProducerQueueMember {
    public static void main( String[] args ) throws Exception {

        HazelcastInstance hz = Hazelcast.newHazelcastInstance();

        IQueue<Integer> queue = hz.getQueue( "queue" );
        for ( int k = 1; k < 100; k++ ) {
            queue.put( k );
            System.out.println( "Producing: " + k );
            Thread.sleep(1000);
        }
        queue.put( -1 );
        System.out.println( "Producer Finished!" );
    }
}
```

```
public class ConsumerQueueMember {
    public static void main( String[] args
    ) throws Exception {
        HazelcastInstance hz =
        Hazelcast.newHazelcastInstance();
        IQueue<Integer> queue =
        hz.getQueue( "queue" );

        while ( true ) {
            int item = queue.take();
            System.out.println( "Consumed: " + item );
            if ( item == -1 ) {
                queue.put( -1 );
                break;
            }
            Thread.sleep( 5000 );
        }
        System.out.println( "Consumer Finished!" );
    }
}
```

# Data Structures : MultiMap

- Hazelcast fournit des implémentations distribuées des structures de données de types :

- Map
- Queue
- **MultiMap**
- Set
- List
- RingBuffer

```
HazelcastInstance hazelcastInstance =  
Hazelcast.newHazelcastInstance();  
MultiMap<String , String > map =  
hazelcastInstance.getMultiMap( "map" );  
map.put( "a", "1" );  
map.put( "a", "2" );  
map.put( "b", "3" );  
System.out.println( "PutMember:Done" );  
  
for (String key: map.keySet()){  
    Collection<String> values = map.get(key);  
    System.out.printf("%s -> %s\n", key, values);  
}
```

```
b → [3]  
a → [2, 1]
```

# Data Structures : Set

- Hazelcast Set does not allow duplicate elements.
- Hazelcast Set does not preserve the order of elements.
- Hazelcast Set is a non-partitioned data structure: all the data that belongs to a set lives on one single partition in that member.
- Hazelcast Set cannot be scaled beyond the capacity of a single machine. Since the whole set lives
- on a single partition, storing a large amount of data on a single set may cause memory pressure.
- Therefore, you should use multiple sets to store a large amount of data. This way, all the sets are spread across the cluster, sharing the load.
- A backup of Hazelcast Set is stored on a partition of another member in the cluster so that data is not lost in the event of a primary member failure.
- All items are copied to the local member and iteration occurs locally.

```
HazelcastInstance hz =  
Hazelcast.newHazelcastInstance();  
  
ISet<String> set = hz.getSet("set");  
  
set.add("Tokyo");  
set.add("Paris");  
set.add("London");  
set.add("New York");  
System.out.println("Putting finished!");  
  
for (String value:set){  
    System.out.println(value);  
}
```



# Data Structures : List

- Hazelcast List (IList) is similar to Hazelcast Set, but it also allows duplicate elements.
  - Besides allowing duplicate elements, Hazelcast List preserves the order of elements.
  - Hazelcast List is a non-partitioned data structure where values and each backup are represented by their own single partition.
  - Hazelcast List cannot be scaled beyond the capacity of a single machine.
  - All items are copied to local and iteration occurs locally.

```
HazelcastInstance hz =  
Hazelcast.newHazelcastInstance();  
  
IList<String> myList = hz.getList("myList");  
  
myList.add("Tokyo");  
myList.add("Paris");  
myList.add("London");  
myList.add("New York");  
System.out.println("Putting finished!");  
  
for (String value:myList){  
    System.out.println(value);  
}
```

# Data Structures : RingBuffer

- Hazelcast Ringbuffer is a replicated but not partitioned data structure that stores its data in a ringlike structure.
- You can think of it as a circular array with a given capacity. Each Ringbuffer has a tail and a head.
- The tail is where the items are added and the head is where the items are overwritten or expired.
- You can reach each element in a Ringbuffer using a sequence ID, which is mapped to the elements between the head and tail (inclusive) of the Ringbuffer.
- Ringbuffer can sometimes be a better alternative than an Hazelcast IQueue.
- Unlike IQueue, Ringbuffer does not remove the items, it only reads items using a certain position.

```
HazelcastInstance hz =  
Hazelcast.newHazelcastInstance();  
  
Ringbuffer<String> ringbuffer =  
hz.getRingbuffer("rb");  
  
ringbuffer.add("Item 1");  
ringbuffer.add("Item 2");  
  
long sequence = ringbuffer.headSequence();  
while(true){  
    String item = ringbuffer.readOne(sequence);  
    sequence++;  
    // process item  
}
```

# Data Structures : Topic

- Hazelcast provides a distribution mechanism for publishing messages that are delivered to multiple subscribers.
- This is also known as a publish/subscribe (pub/sub) messaging model.
- Publishing and subscribing operations are cluster wide.
- When a member subscribes to a topic, it is actually registering for messages published by any member in the cluster, including the new members that joined after you add the listener.

```
public class TopicPublisher {  
    public static void main(String[] args) {  
        HazelcastInstance hz = Hazelcast.newHazelcastInstance();  
        ITopic<String> topic = hz.getTopic("topic");  
        topic.publish("My message : Hi...");  
    }  
}
```

```
public class TopicSubscriber {  
    public static void main(String[] args) {  
        HazelcastInstance hz = Hazelcast.newHazelcastInstance();  
        ITopic<String> topic = hz.getTopic("topic");  
        topic.addMessageListener(new MessageListener<String>() {  
            @Override  
            public void onMessage(Message<String> message) {  
                System.out.println("Received message  
:" + message.getMessageObject());  
            }  
        });  
        System.out.println("Subscribed");  
    }  
}
```

# Data Structures : Topic

- Hazelcast provides a distribution mechanism for publishing messages that are delivered to multiple subscribers.
- This is also known as a publish/subscribe (pub/sub) messaging model.
- Publishing and subscribing operations are cluster wide.
- When a member subscribes to a topic, it is actually registering for messages published by any member in the cluster, including the new members that joined after you add the listener.

```
public class TopicPublisher {  
    public static void main(String[] args) {  
        HazelcastInstance hz = Hazelcast.newHazelcastInstance();  
        ITopic<String> topic = hz.getTopic("topic");  
        topic.publish("My message : Hi...");  
    }  
}
```

```
public class TopicSubscriber {  
    public static void main(String[] args) {  
        HazelcastInstance hz = Hazelcast.newHazelcastInstance();  
        ITopic<String> topic = hz.getTopic("topic");  
        topic.addMessageListener(new MessageListener<String>() {  
            @Override  
            public void onMessage(Message<String> message) {  
                System.out.println("Received message  
:" + message.getMessageObject());  
            }  
        });  
        System.out.println("Subscribed");  
    }  
}
```

# Data Structures : Configuring a Topic

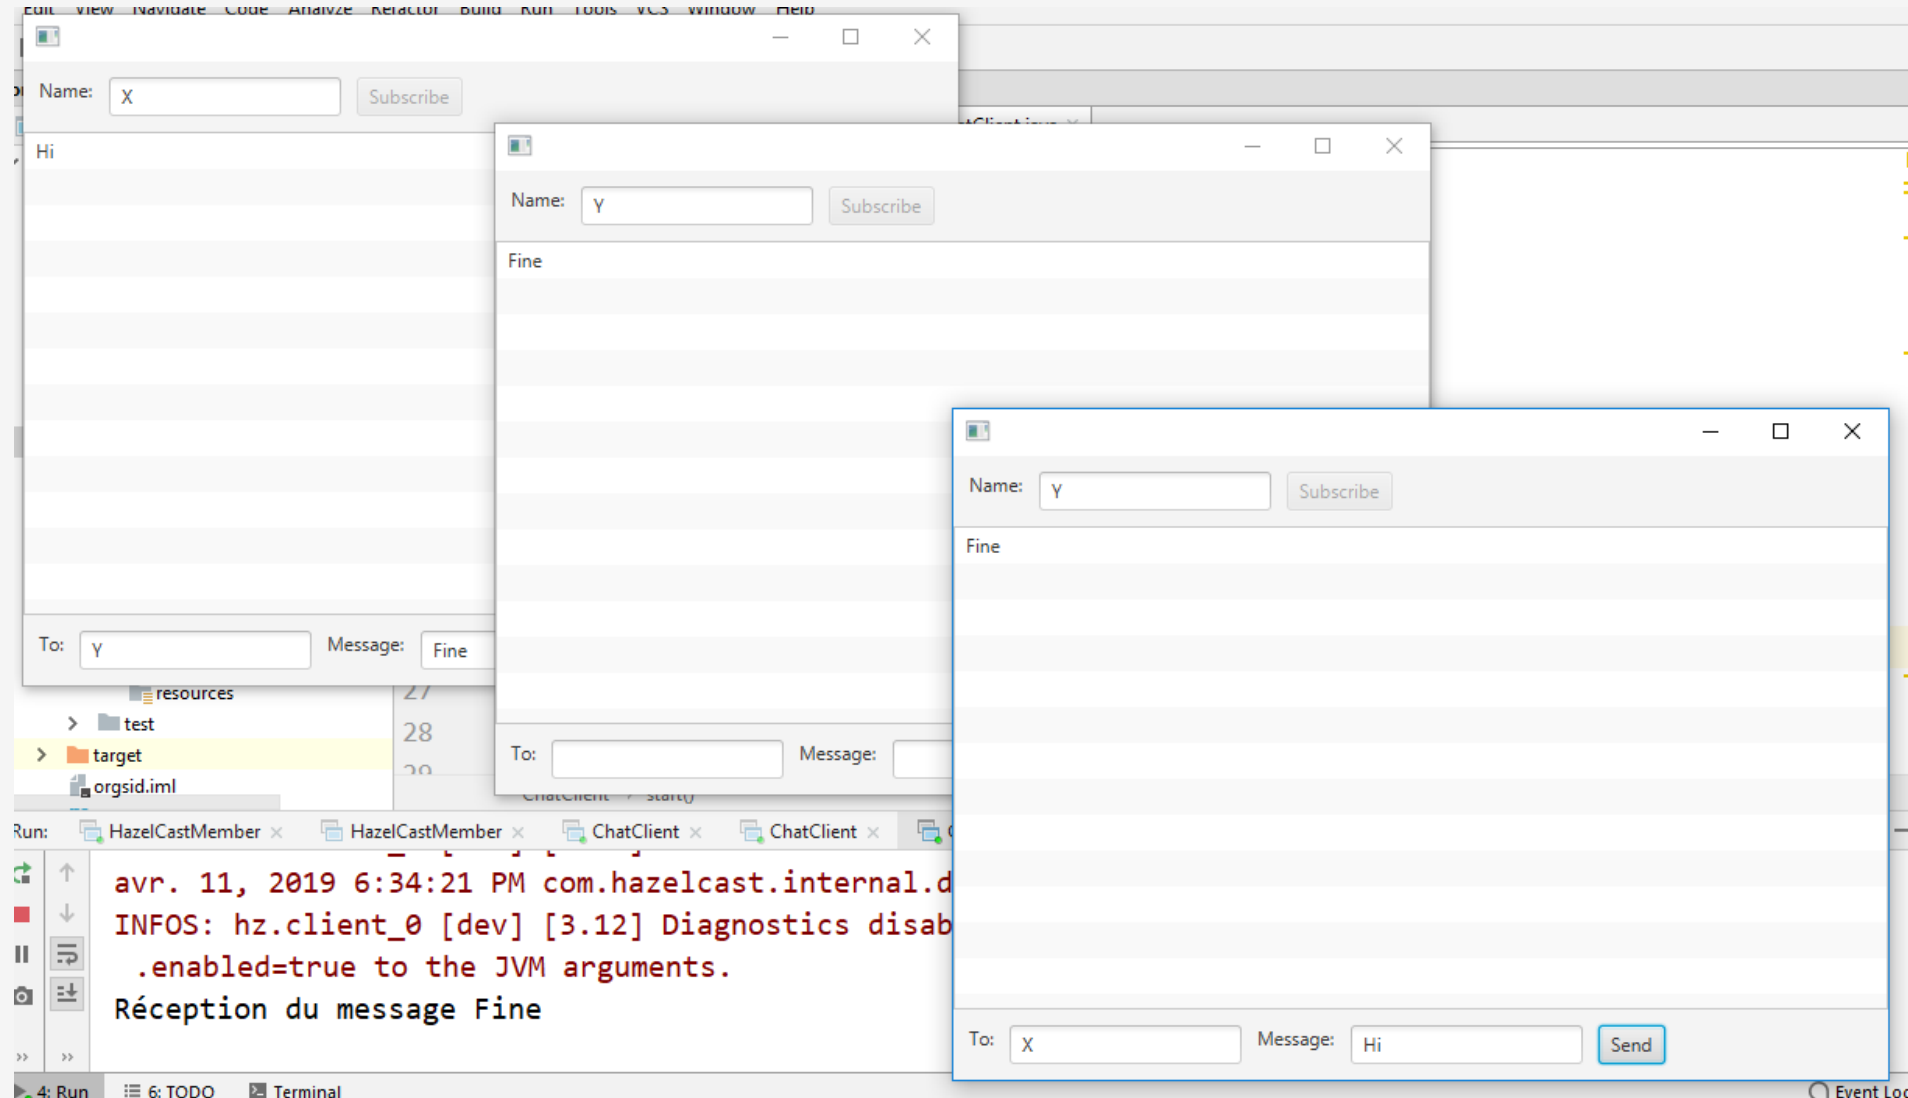
```
import com.hazelcast.config.Config; import com.hazelcast.config.ListenerConfig;
import com.hazelcast.config.ManagementCenterConfig; import com.hazelcast.config.TopicConfig;
import com.hazelcast.core.*;
public class HazelcastNode {
    public static void main(String[] args) {
        ManagementCenterConfig managementCenterConfig=new ManagementCenterConfig();
        managementCenterConfig.setEnabled(true);
        managementCenterConfig.setUrl("http://localhost:8080/hazelcast-mancenter");

        Config cfg = new Config();
        cfg.setManagementCenterConfig(managementCenterConfig);

        TopicConfig topicConfig=new TopicConfig();
        topicConfig.setName("topic");
        topicConfig.setGlobalOrderingEnabled(true);
        topicConfig.setStatisticsEnabled(true);
        topicConfig.addMessageListenerConfig(new ListenerConfig(new MessageListener<String>() {
            @Override
            public void onMessage(Message<String> message) {
                System.out.println("Message :"+message.getMessageObject());
            }
        }));
        cfg.addTopicConfig(topicConfig);
        HazelcastInstance instance = Hazelcast.newHazelcastInstance(cfg);
    }
}
```



# Data Structures : Client Chat avec Hazelcast Topic



# Data Structures : Client Chat avec Hazelcast Topic

```
import com.hazelcast.client.HazelcastClient;import com.hazelcast.client.impl.clientside.HazelcastClientInstanceImpl;
import com.hazelcast.core.Hazelcast; import com.hazelcast.core.HazelcastInstance;
import com.hazelcast.core.ITopic; import javafx.application.Application; import javafx.collections.FXCollections;
import javafx.collections.ObservableArray; import javafx.collections.ObservableList;import javafx.geometry.Insets;
import javafx.scene.Scene;import javafx.scene.control.*; import javafx.scene.layout.BorderPane;
import javafx.scene.layout.HBox; import javafx.stage.Stage; import java.util.ArrayList; import java.util.List;
public class ChatClient extends Application {
    private HazelcastInstance hazelcastInstance;
    @Override
    public void start(Stage primaryStage) throws Exception {
        hazelcastInstance= HazelcastClient.newHazelcastClient();
        BorderPane borderPane=new BorderPane();
        Label labelName=new Label("Name:");
        TextField textFieldName=new TextField();
        Button buttonSubscribe=new Button("Subscribe");
        Label labelTo=new Label("To:");
        TextField textFieldTo=new TextField();
        Button buttonSend=new Button("Send");
        Label labelMessage=new Label("Message:");
        TextField textFieldMessage=new TextField();
        List<String> messages=new ArrayList<>();
        ObservableList<String> observableList= FXCollections.observableList(messages);
        ListView<String> listView=new ListView<>(observableList);
```

# Data Structures : Client Chat avec Hazelcast Topic

```
HBox hBoxTop=new HBox(10); hBoxTop.setPadding(new Insets(10));
    hBoxTop.getChildren().addAll(labelName,textFieldName,buttonSubscribe);
    HBox hBoxBottom=new HBox(10);hBoxBottom.setPadding(new Insets(10));
    hBoxBottom.getChildren().addAll(labelTo,textFieldTo,labelMessage,textFieldMessage,buttonSend);
    borderPane.setTop(hBoxTop);
    borderPane.setBottom(hBoxBottom);
    borderPane.setCenter(listView);
    Scene scene=new Scene(borderPane,600,400);
    primaryStage.setScene(scene);
    primaryStage.show();

    buttonSubscribe.setOnAction(evt->{
        ITopic<String> topic=hazelcastInstance.getTopic(textFieldName.getText());
        topic.addListener(message->{
            System.out.println("Réception du message "+message.getMessageObject());
            observableList.add(message.getMessageObject());
        });
        buttonSubscribe.setDisable(true);
    });

    buttonSend.setOnAction(evt->{
        ITopic topicTo=hazelcastInstance.getTopic(textFieldTo.getText());
        topicTo.publish(textFieldMessage.getText());
    });
}
```

# Loading and Storing Persistent Data

---

- Hazelcast allows you to load and store the distributed Structures entries from/to a persistent data store such as a relational database.
- To do this, for example in Map Structure, you can use Hazelcast's MapStore and MapLoader interfaces.
- When you provide a MapLoader implementation and request an entry (IMap.get()) that does not exist in memory, MapLoader's load method loads that entry from the data store.
- This loaded entry is placed into the map and will stay there until it is removed or evicted.
- When a MapStore implementation is provided, an entry is also put into a user defined data store.
- Data store needs to be a centralized system that is accessible from all Hazelcast members. Persistence to a local file system is not supported.

# Loading and Storing Persistent Data

```
package a;
import com.hazelcast.core.MapStore; import java.util.Collection;
import java.util.Map;
public class PersonMapStore implements MapStore<String,Person> {
    @Override
    public void store(String s, Person person) {
        System.out.println("*****");
        System.out.println("Storing Person " + person.getName());
    }
    @Override
    public void storeAll(Map<String, Person> map) { }
    @Override
    public void delete(String s) { }
    @Override
    public void deleteAll(Collection<String> collection) { }
    @Override
    public Person load(String s) { return null; }
    @Override
    public Map<String, Person> loadAll(Collection<String> collection) {
        return null;
    }
    @Override
    public Iterable<String> loadAllKeys() {
        return null;
    }
}
```

```
package a;
import java.io.Serializable;
public class Person implements Serializable {
    private Long id;
    private String name;

    public Person(Long id, String name) {
        this.id = id;
        this.name = name;
    }

    public Long getId() {
        return id;
    }

    public void setId(Long id) {
        this.id = id;
    }

    public String getName() {
        return name;
    }

    public void setName(String name) {
        this.name = name;
    }
}
```

# Loading and Storing Persistent Data

```
Config config = new Config();
MapConfig mapConfig = config.getMapConfig("default");
mapConfig.setName("MyQueue")
    .setBackupCount(1)
    .setStatisticsEnabled(true);

mapConfig.getMapStoreConfig()
    .setEnabled(true)
    .setClassName("a.PersonMapStore")
    .setProperty("binary", "false");
config.addMapConfig(mapConfig);
HazelcastInstance hz=Hazelcast.newHazelcastInstance(config);

Map<String, Person> capitalcities = hz.getMap( "MyQueue" );
capitalcities.put( "1", new Person(1L,"Mohamed") );
capitalcities.put( "1", new Person(2L,"Imane") );
```

# Data Structures

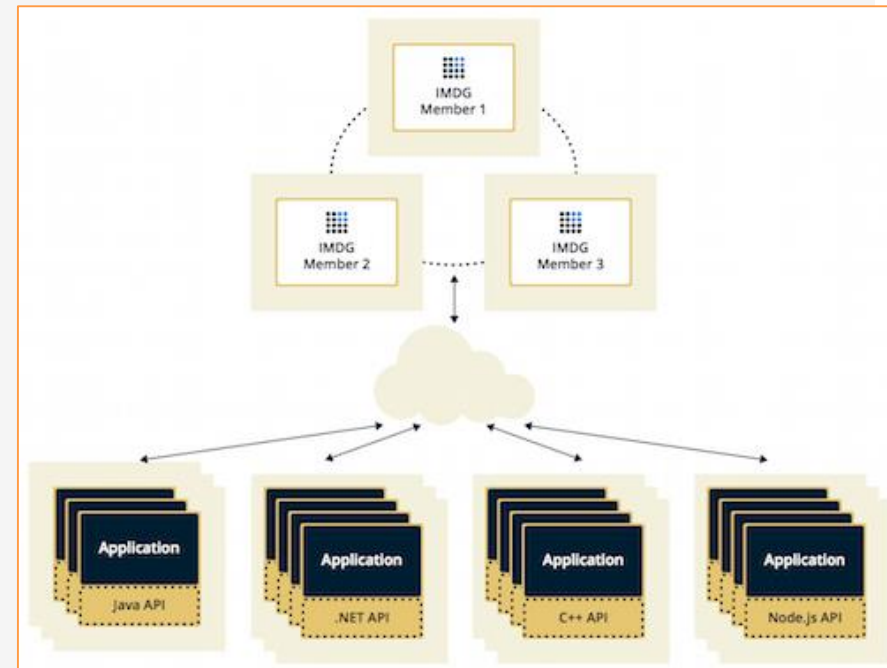
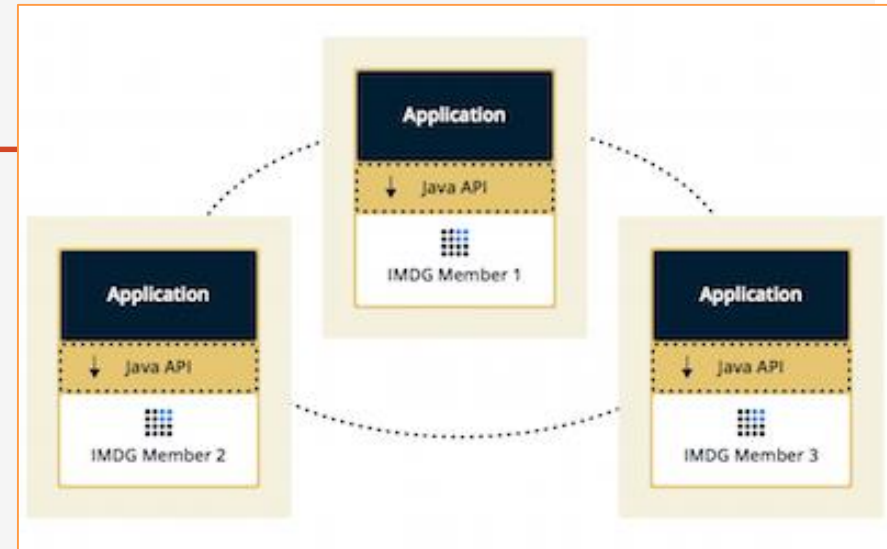
---

- Hazelcast fournit des implémentations distribuées des structures de données de types :
  - Map
  - Queue
  - MultiMap
  - Set
  - List
  - RingBuffer



# Distributed Computing with Hazelcast

- Hazelcast offers IExecutorService for you to use in distributed environments.
- It implements java.util.concurrent.ExecutorService to serve the applications requiring computational and data processing power.
- With IExecutorService, you can execute tasks asynchronously and perform other useful tasks.
- If your task execution takes longer than expected, you can cancel the task execution.
- Tasks should be Serializable since they are distributed.
- In the Java Executor framework, you implement tasks two ways: Callable or Runnable.
  - Callable: If you need to return a value and submit it to Executor, implement the task as **java.util.concurrent.Callable**.
  - Runnable: If you do not need to return a value, implement the task as **java.util.concurrent.Runnable**.



# Distributed Computing with Hazelcast : Callable Task

```
import com.hazelcast.core.HazelcastInstance;
import com.hazelcast.core.HazelcastInstanceAware;
import com.hazelcast.core.IMap;


import java.io.Serializable;
import java.util.concurrent.Callable;

public class SumTask implements Callable<Integer>, Serializable, HazelcastInstanceAware {
    private transient HazelcastInstance hazelcastInstance;
    @Override
    public void setHazelcastInstance(HazelcastInstance hazelcastInstance) {
        this.hazelcastInstance=hazelcastInstance;
    }
    @Override
    public Integer call() throws Exception {
        IMap<Integer, Integer> map = hazelcastInstance.getMap( "inputData" );
        int result = 0;
        for ( Integer key : map.localKeySet() ) {
            System.out.println( "Calculating for key: " + key );
            result += map.get( key );
        }
        System.out.println( "Local Result: " + result );
        return result;
    }
}
```

# Distributed Computing with Hazelcast : Callable Task


```
import com.hazelcast.config.Config;
import com.hazelcast.config.ManagementCenterConfig;
import com.hazelcast.core.Hazelcast;
import com.hazelcast.core.HazelcastInstance;

public class HazelCastWorker {
    public static void main(String[] args) {
        Config config=new Config();
        ManagementCenterConfig centerConfig=new
ManagementCenterConfig();
        centerConfig.setEnabled(true);
        centerConfig.setUrl("http://localhost:8080/hazelcast-
mancenter");
        config.setManagementCenterConfig(centerConfig);
        HazelcastInstance hazelcastInstance=
Hazelcast.newHazelcastInstance();
    }
}
```



# Distributed Computing with Hazelcast : Callable Task

```
import com.hazelcast.client.HazelcastClient; import com.hazelcast.core.HazelcastInstance;
import com.hazelcast.core.IExecutorService; import com.hazelcast.core.IMap; import com.hazelcast.core.Member;
import java.util.Map; import java.util.concurrent.ExecutionException; import java.util.concurrent.Future;
public class TaskProducer {
    public static void main(String[] args) throws ExecutionException, InterruptedException {
        HazelcastInstance hazelcastInstance= HazelcastClient.newHazelcastClient();
        IMap<Integer,Integer> data=hazelcastInstance.getMap("inputData");
        for (int i = 0; i <10000 ; i++) { data.put(i,1); }
        IExecutorService executorService=hazelcastInstance.getExecutorService("default");
        /*
        Future<Integer> response=executorService.submit(new SumTask());
        System.out.println("Result="+response.get());
        */
        Map<Member,Future<Integer>> response=
        executorService.submitToAllMembers(new SumTask());
        double reduceSum=0;
        for (Member member:response.keySet()){
            System.out.println("*****");
            System.out.println(member.getAddress());
            System.out.println(response.get(member).get());
            reduceSum+=response.get(member).get();
            System.out.println("*****");
        }
        System.out.println("Total SUM="+reduceSum);
    }
}
```



```
*****
[127.0.0.1]:5703
3319
*****
*****
[127.0.0.1]:5701
3355
*****
*****
[127.0.0.1]:5702
3326
*****
Total SUM=10000.0
```

# Architecture de Hazelcast

